



CS6203: Advanced Topics in Database Management Systems

The Case for Learned Index Structures

HUANG liu

e0575772@u.nus.edu

Introduction

Range Index

Point Index

Existence Index

Related Work & Conclusion



Introduction

Index structures

DATA ACCESS PATTERNS

range request
(from sorted array)

single key look-ups
(from unsorted array)

check for record
existence



B-tree



HashMap



Bloom Filter



General cases



Specific case

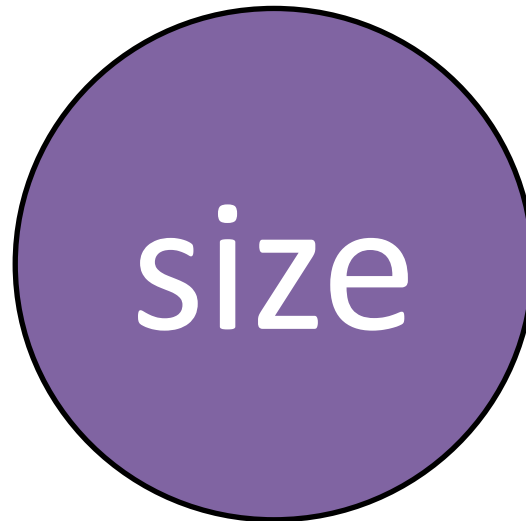
Array: [100, 101, ... 100M]

B-Tree: $O(\log(n))$

Array[key - 100]: $O(1)$

With the knowledge of data distribution, we can highly optimize the index structure.

Additional notes
to lectures from data
Data structures using ML



Learned Index



+

Tradition



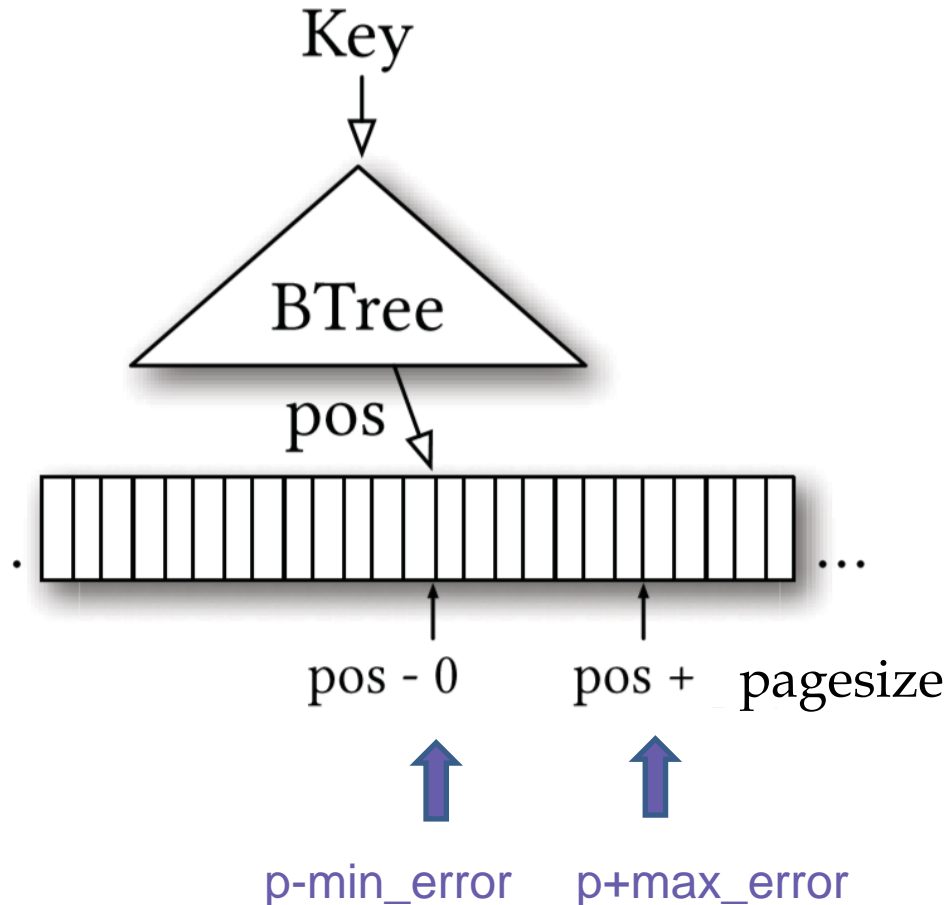
~~Tradition~~



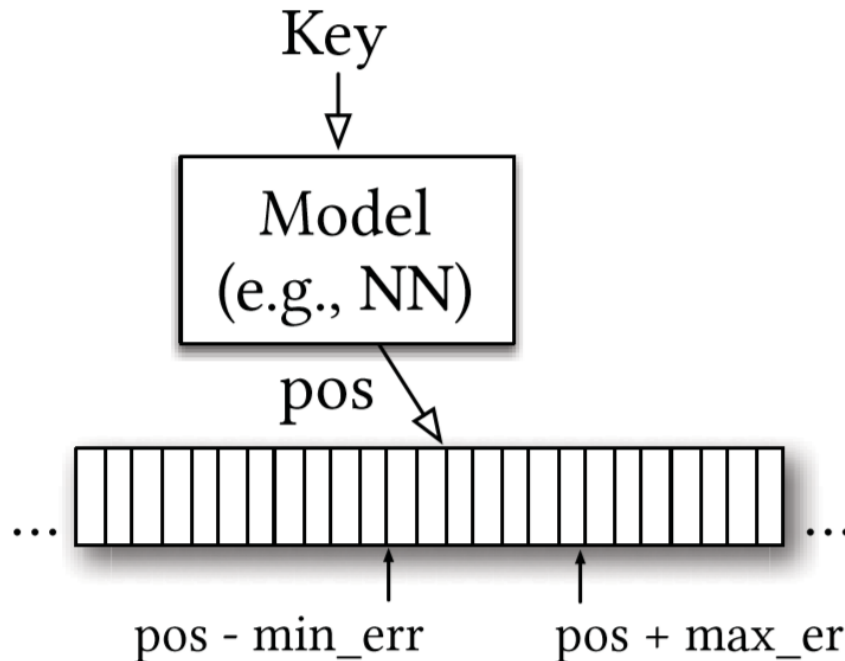
Range Index

B-Tree is a model

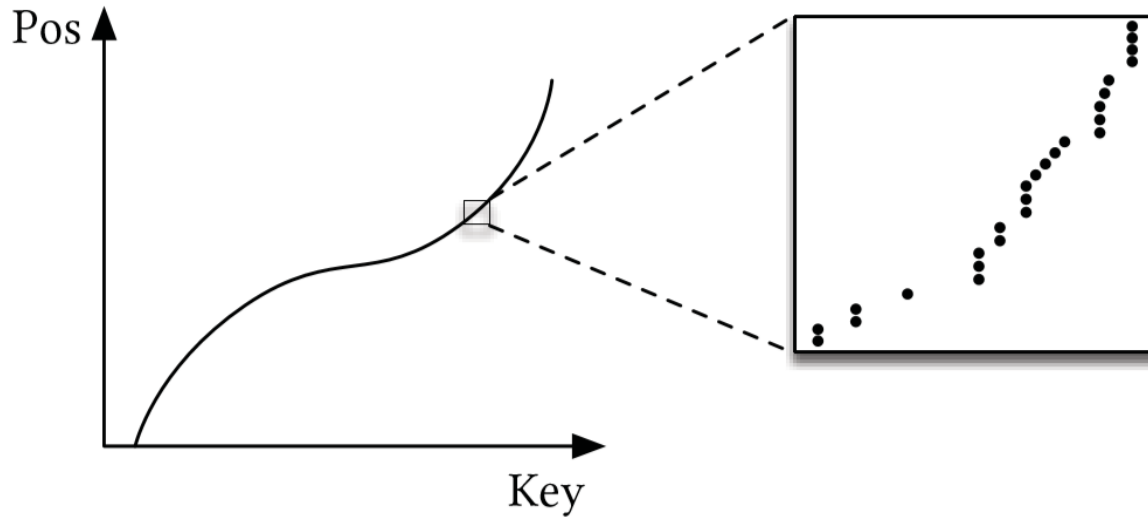
Assumptions:
In-memory
database
(i.e., read-only)
No insert



Another view of index structure



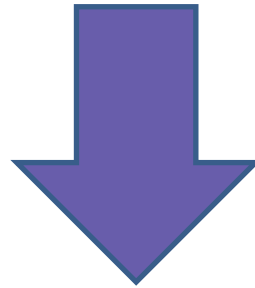
Min_err and Max_err are known from the training process.
Eg. For B-trees, the error is the page_size.



$$\text{pos} = F(\text{key}) * \# \text{Keys}$$
$$F(\text{key}) = P(x \leq \text{key})$$

- A model to predict the position of a given key is actually building a model to present the **CDF(Cumulative Distribution Function)** of data.

B-tree: regression tree



ML models: linear regression, NN...

Benefits of learned index models

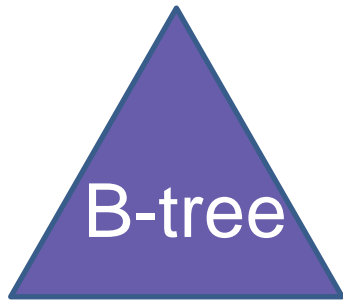
- **Smaller index**: less main-memory storage.
- **Faster lookup**: eg. Linear function.
- **More parallelism**: use multiplications instead of if-statement.

Benefits of learned index depend on:

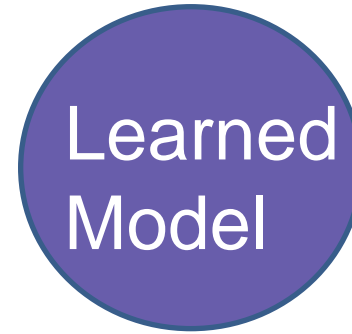
- How accurately the model represents the observed CDF
- Architecture, implementation details, payload

A First Naive Learned Index

- Tensorflow
- 200M web-server log records
- Input features: timestamps
- Labels: positions in sorted array
- 2 fc , 32 neurons/layer, ReLu activated



300ns

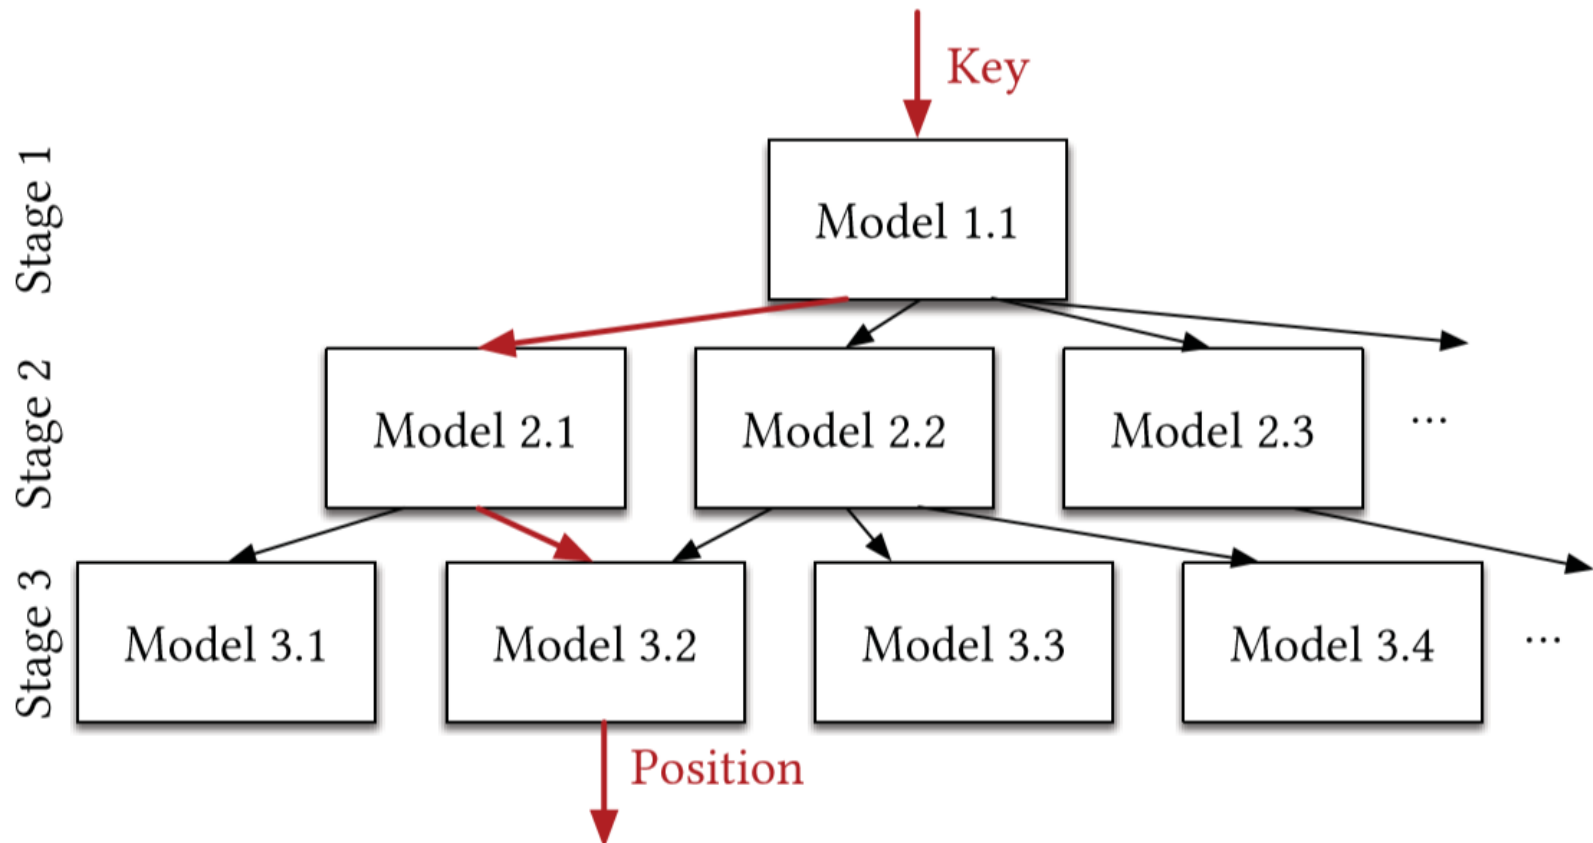


80000ns

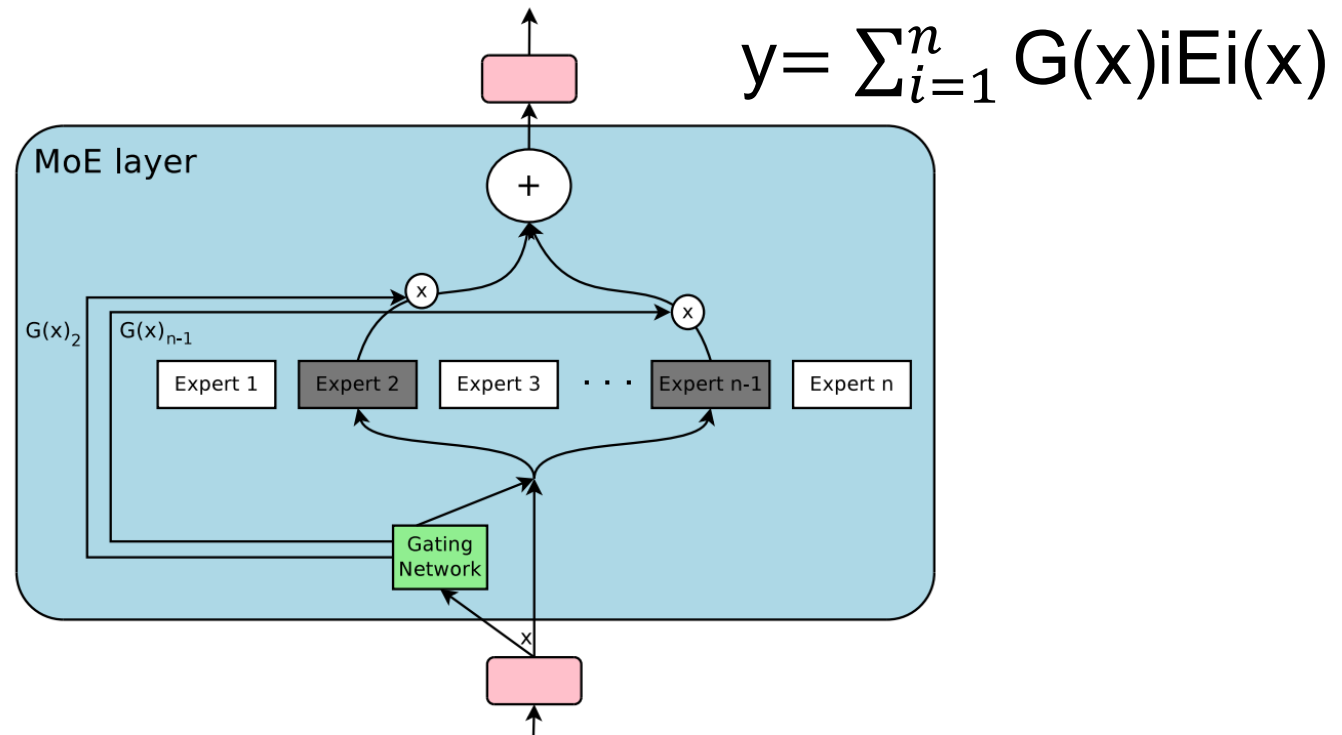
REASONS:

1. Tensorflow is designed for large models
2. B-tree can easily overfit the training data and can predict the position precisely
3. B-trees are cache efficient

Solution: Recursive Model Index

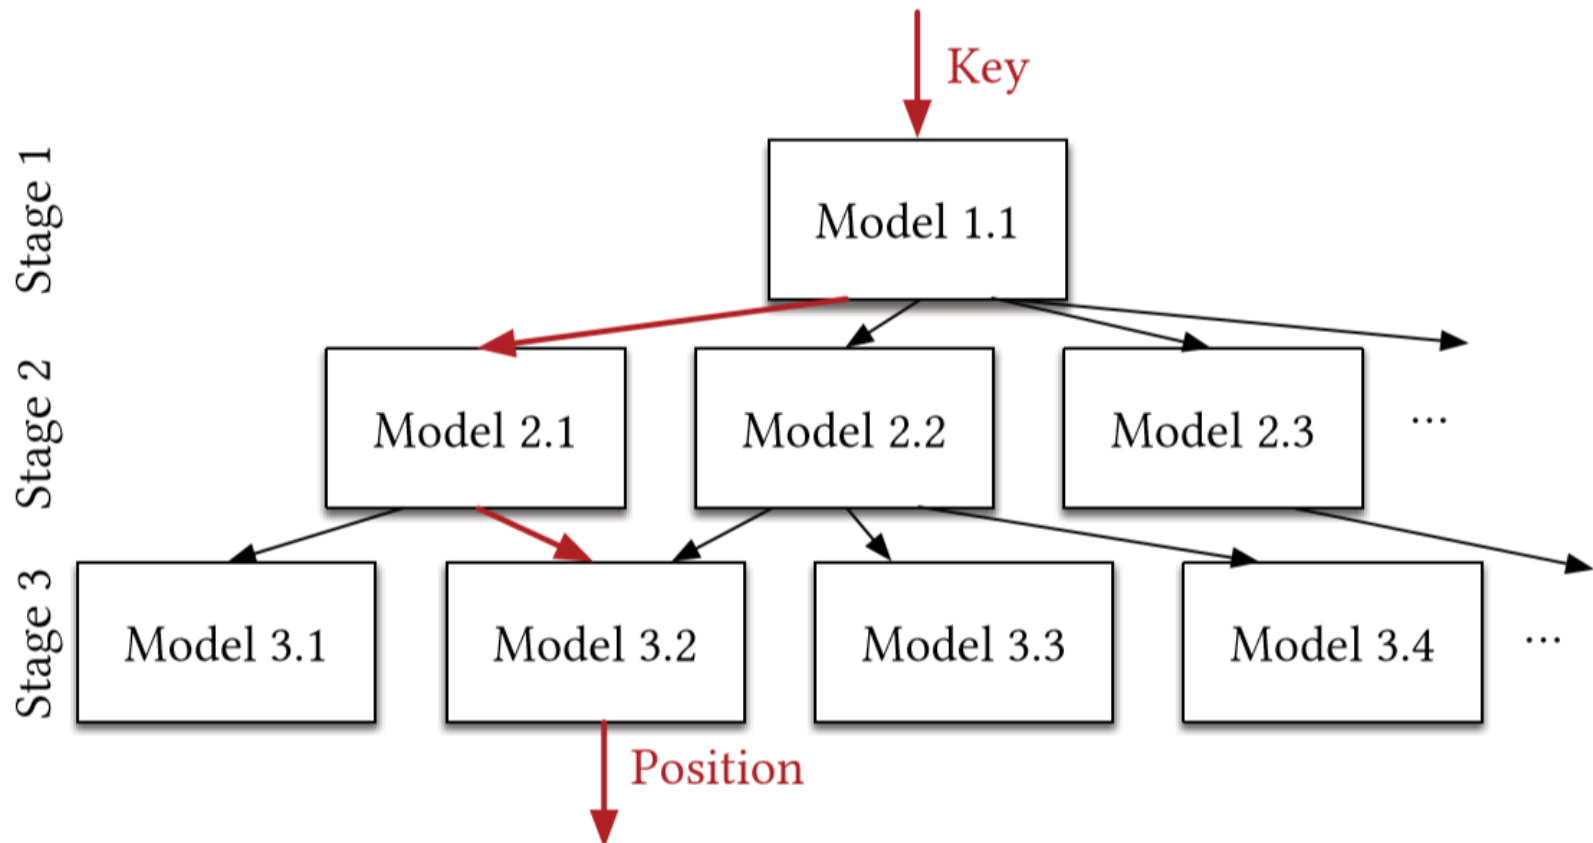


A Mixture of Experts (MoE) layer



N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture of experts layer. arXiv preprint arXiv:1701.06538, 2017.

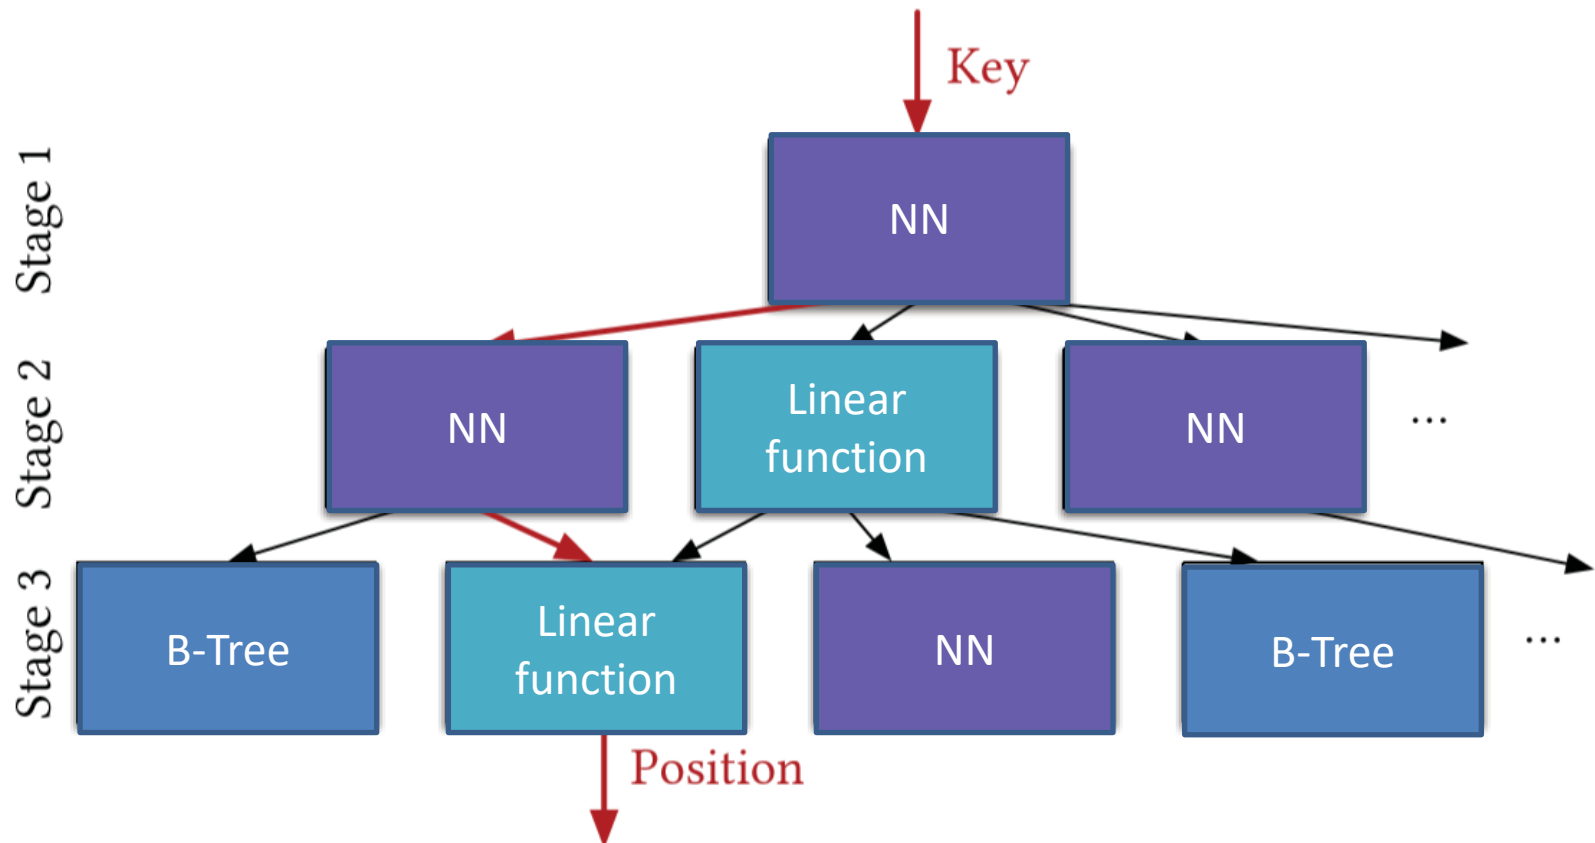
Solution: Recursive Model Index



Advantages of RMI

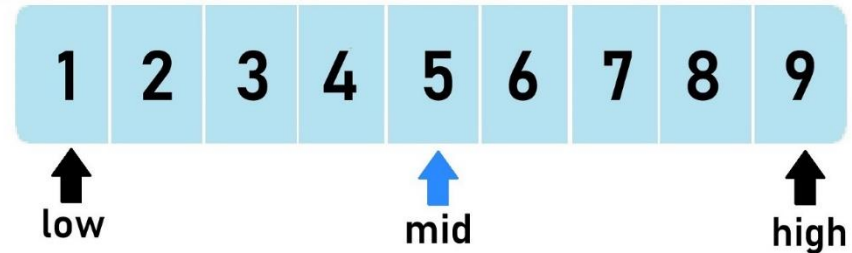
- separates model size and complexity from execution cost
- easy to learn the overall data distribution
- easier to solve the last-mile accuracy problem.
- no search process is required in-between stages.

Hybrid Indexes



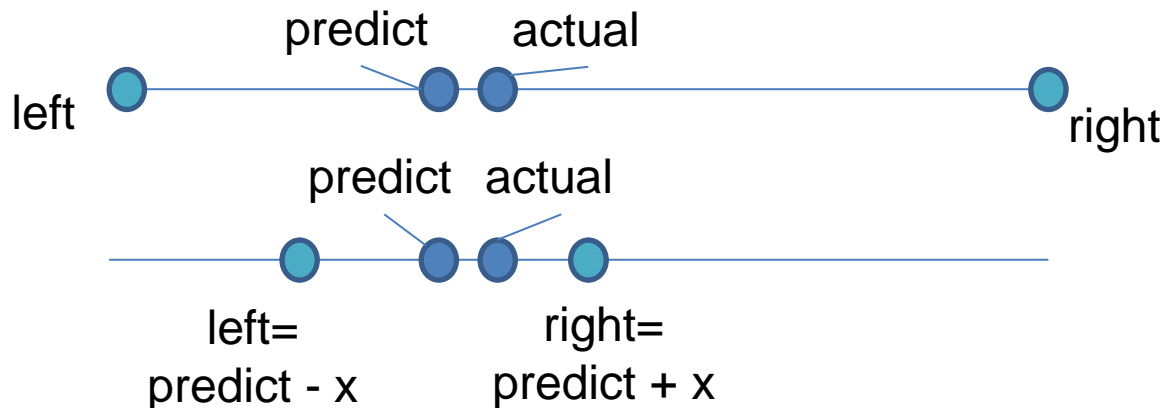
Search Strategies

Binary Search



<https://images.app.goo.gl/P7mvfhx1vaXCKA69>

Quaternary Search



<https://huangliu0909.github.io>

Indexing Strings

- Tokenization: $x = (x_1, x_2, \dots, x_n)$
- x_i is the ASCII decimal value of the corresponding character.
- Set a fixed length N : For strings with length $n < N$, we set $x_i = 0$ for $i > n$.

Result

Type	Config	Map Data			Web Data			Log-Normal Data		
		Size (MB)	Lookup (ns)	Model (ns)	Size (MB)	Lookup (ns)	Model (ns)	Size (MB)	Lookup (ns)	Model (ns)
Btree	page size: 32	52.45 (4.00x)	274 (0.97x)	198 (72.3%)	51.93 (4.00x)	276 (0.94x)	201 (72.7%)	49.83 (4.00x)	274 (0.96x)	198 (72.1%)
	page size: 64	26.23 (2.00x)	277 (0.96x)	172 (62.0%)	25.97 (2.00x)	274 (0.95x)	171 (62.4%)	24.92 (2.00x)	274 (0.96x)	169 (61.7%)
	page size: 128	13.11 (1.00x)	265 (1.00x)	134 (50.8%)	12.98 (1.00x)	260 (1.00x)	132 (50.8%)	12.46 (1.00x)	263 (1.00x)	131 (50.0%)
	page size: 256	6.56 (0.50x)	267 (0.99x)	114 (42.7%)	6.49 (0.50x)	266 (0.98x)	114 (42.9%)	6.23 (0.50x)	271 (0.97x)	117 (43.2%)
	page size: 512	3.28 (0.25x)	286 (0.93x)	101 (35.3%)	3.25 (0.25x)	291 (0.89x)	100 (34.3%)	3.11 (0.25x)	293 (0.90x)	101 (34.5%)
Learned Index	2nd stage models: 10k	0.15 (0.01x)	98 (2.70x)	31 (31.6%)	0.15 (0.01x)	222 (1.17x)	29 (13.1%)	0.15 (0.01x)	178 (1.47x)	26 (14.6%)
	2nd stage models: 50k	0.76 (0.06x)	85 (3.11x)	39 (45.9%)	0.76 (0.06x)	162 (1.60x)	36 (22.2%)	0.76 (0.06x)	162 (1.62x)	35 (21.6%)
	2nd stage models: 100k	1.53 (0.12x)	82 (3.21x)	41 (50.2%)	1.53 (0.12x)	144 (1.81x)	39 (26.9%)	1.53 (0.12x)	152 (1.73x)	36 (23.7%)
	2nd stage models: 200k	3.05 (0.23x)	86 (3.08x)	50 (58.1%)	3.05 (0.24x)	126 (2.07x)	41 (32.5%)	3.05 (0.24x)	146 (1.79x)	40 (27.6%)

- Baseline: B-tree, dense pages
- RMI: 2-stage, NN with 0-2 hidden layers, layer width 4-32

Other baseline

	Lookup Table w/ AVX search	FAST	Fixe-Size Btree w/ interpol. search	Multivariate Learned Index
Time	199 ns	189 ns	280 ns	105 ns
Size	16.3 MB	1024 MB	1.5 MB	1.5 MB

Log normal data with a payload of an eight-byte pointer

Index over String

	Config	Size(MB)	Lookup (ns)	Model (ns)
Btree	page size: 32	13.11 (4.00x)	1247 (1.03x)	643 (52%)
	page size: 64	6.56 (2.00x)	1280 (1.01x)	500 (39%)
	page size: 128	3.28 (1.00x)	1288 (1.00x)	377 (29%)
	page size: 256	1.64 (0.50x)	1398 (0.92x)	330 (24%)
Learned Index	1 hidden layer	1.22 (0.37x)	1605 (0.80x)	503 (31%)
	2 hidden layers	2.26 (0.69x)	1660 (0.78x)	598 (36%)
Hybrid Index	t=128, 1 hidden layer	1.67 (0.51x)	1397 (0.92x)	472 (34%)
	t=128, 2 hidden layers	2.33 (0.71x)	1620 (0.80x)	591 (36%)
	t= 64, 1 hidden layer	2.50 (0.76x)	1220 (1.06x)	440 (36%)
	t= 64, 2 hidden layers	2.79 (0.85x)	1447 (0.89x)	556 (38%)
Learned QS	1 hidden layer	1.22 (0.37x)	1155 (1.12x)	496 (43%)

Non-hybrid RMI with quaternary search

All RMI indexes used 10000 models on the 2nd stage

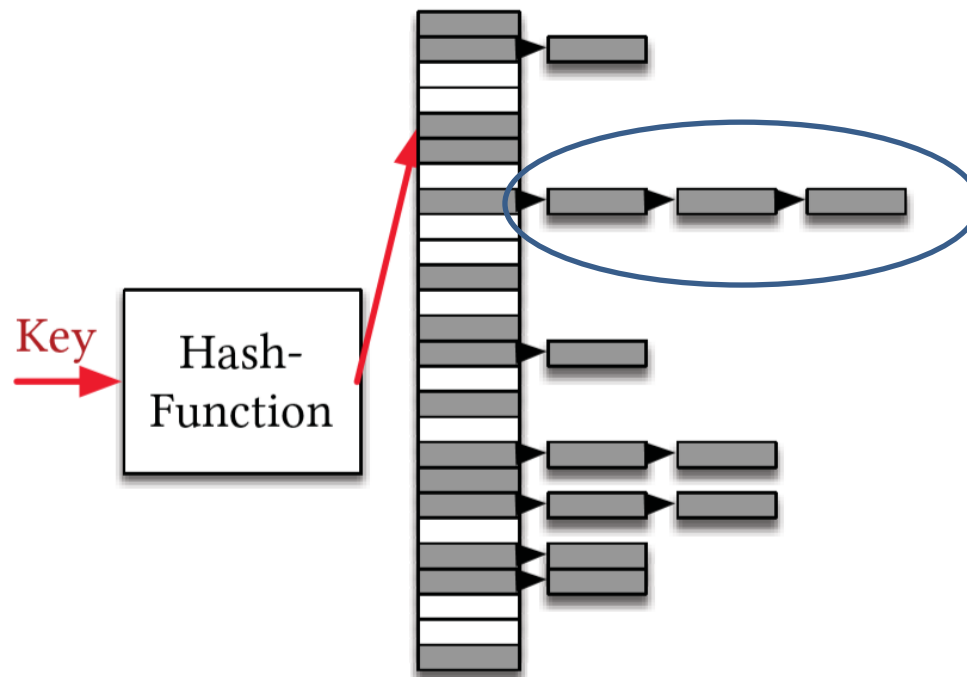


Point Index

Point Index

Point Index is the structure that maps keys to positions inside an array, i.e., Hash-Map

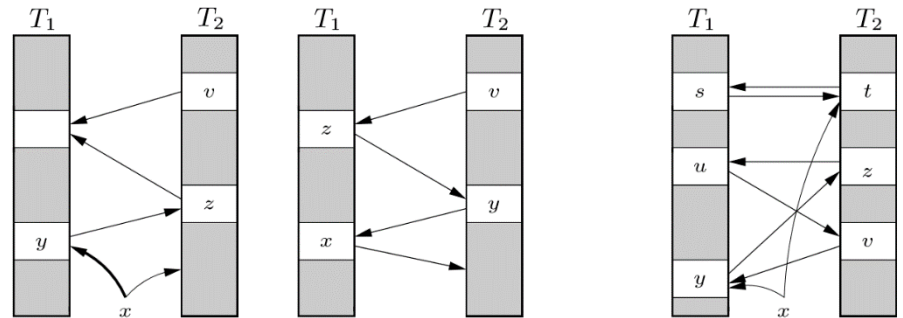
(a) Traditional Hash-Map



Reduce Conflicts

- More than one hashing

R. Pagh and F. F. Rodler. Cuckoo hashing. Journal of Algorithms, 51(2):122– 144, 2004.

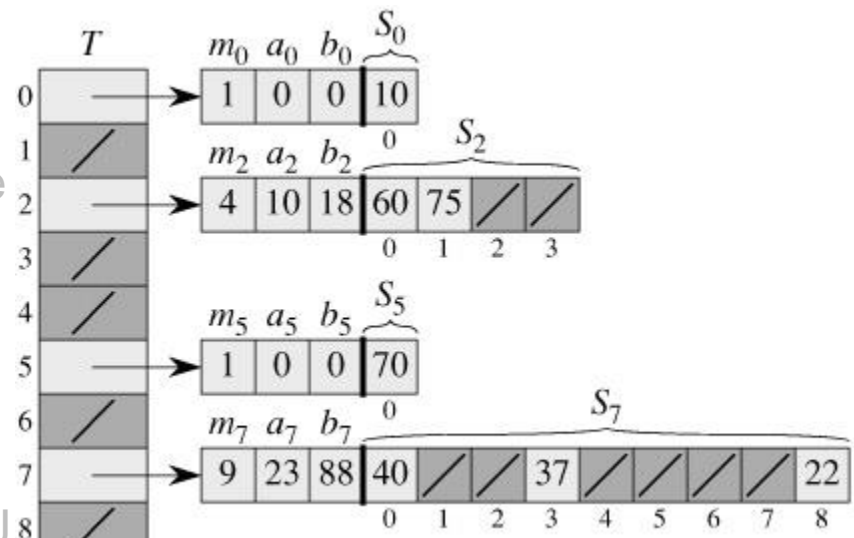


- Perfect hashing

M. Dietzfelbinger, A. Karlin, K. Mehlhorn, F. Meyer auf der Heide, H. Rohnert, and R. E. Tarjan. Dynamic perfect hashing: Upper and lower bounds. SIAM Journal on Computing, 23(4):738–761, 1994.

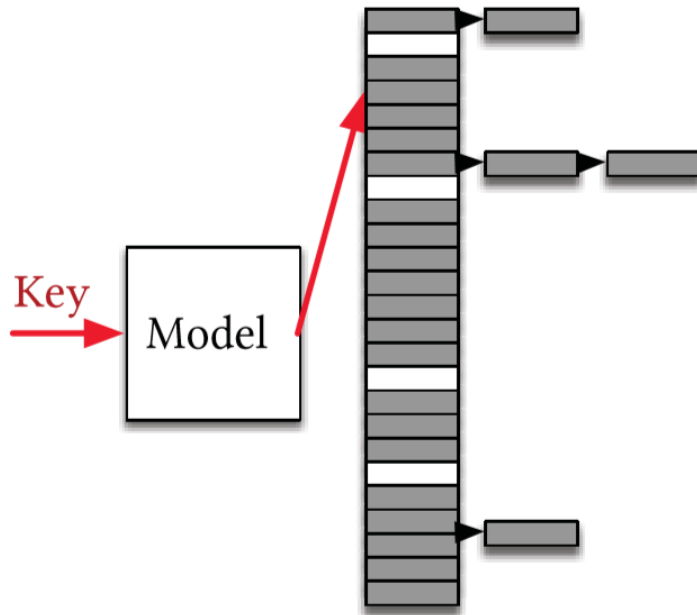
Image from:

<https://images.app.goo.gl/Sk7ADonEwLi6J8E6U7>
<https://huangliu0909.github.io>



Learned hash function

(b) Learned Hash-Map



- $h(k) = F(k) * M$
- M : target size
- Use the function F to learn the empirical CDF of the keys
- Can be combined with chaining or other hash-map type

Result

	% Conflicts Hash Map	% Conflicts Model	Reduction
Map Data	35.3%	07.9%	77.5%
Web Data	35.3%	24.7%	30.0%
Log Normal	35.4%	25.9%	26.7%

Figure 8: Reduction of Conflicts

- Model: 2-stage RMI models with 100k models on the 2nd stage, no hidden layers.
- Baseline: MurmurHash3-like hash-function.

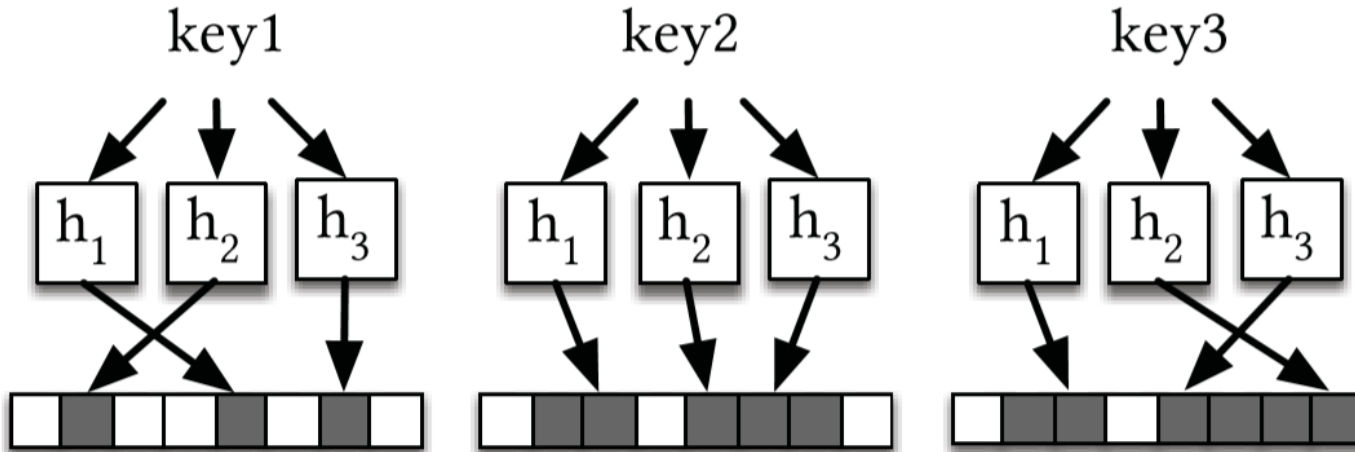
Fewer conflicts lead to fewer cache misses and better performance.



Existence Index

Traditional Bloom-Filter

Insert



Check



key1



key4



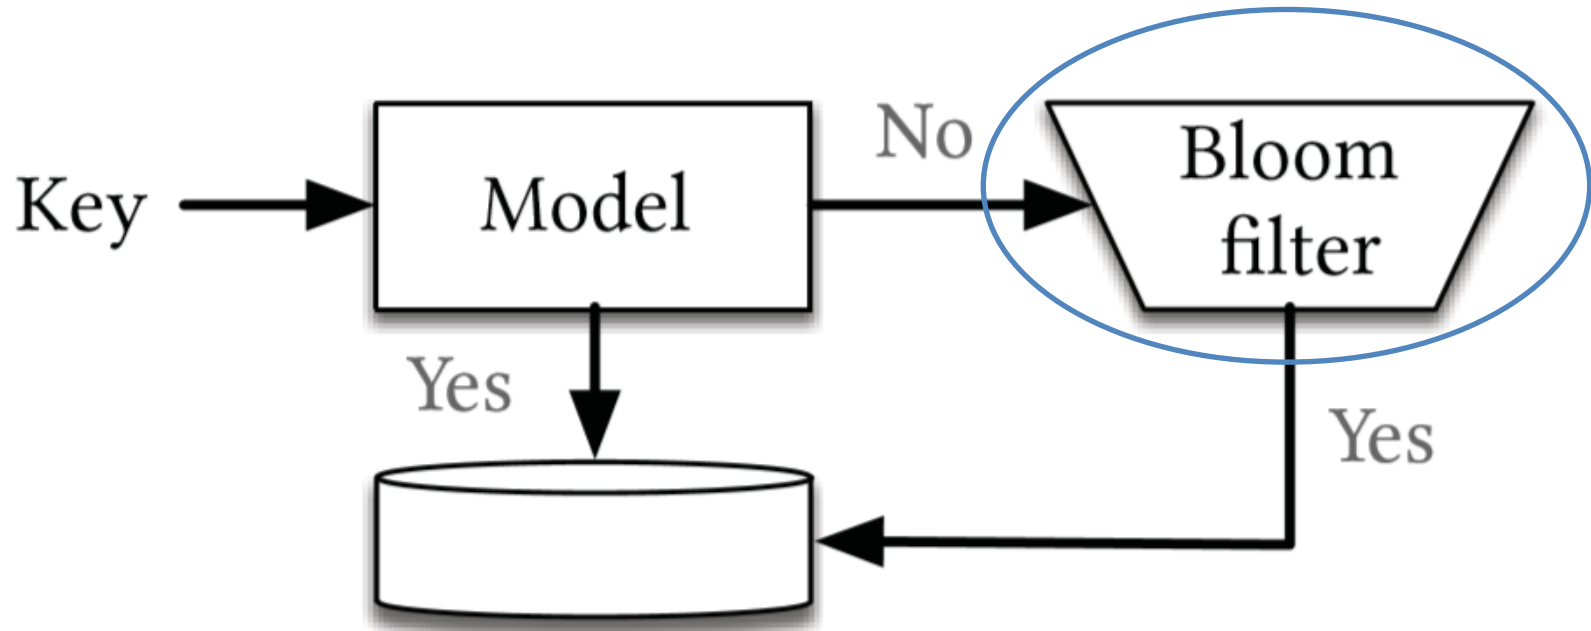
key5



No False Negative

False
Positive !

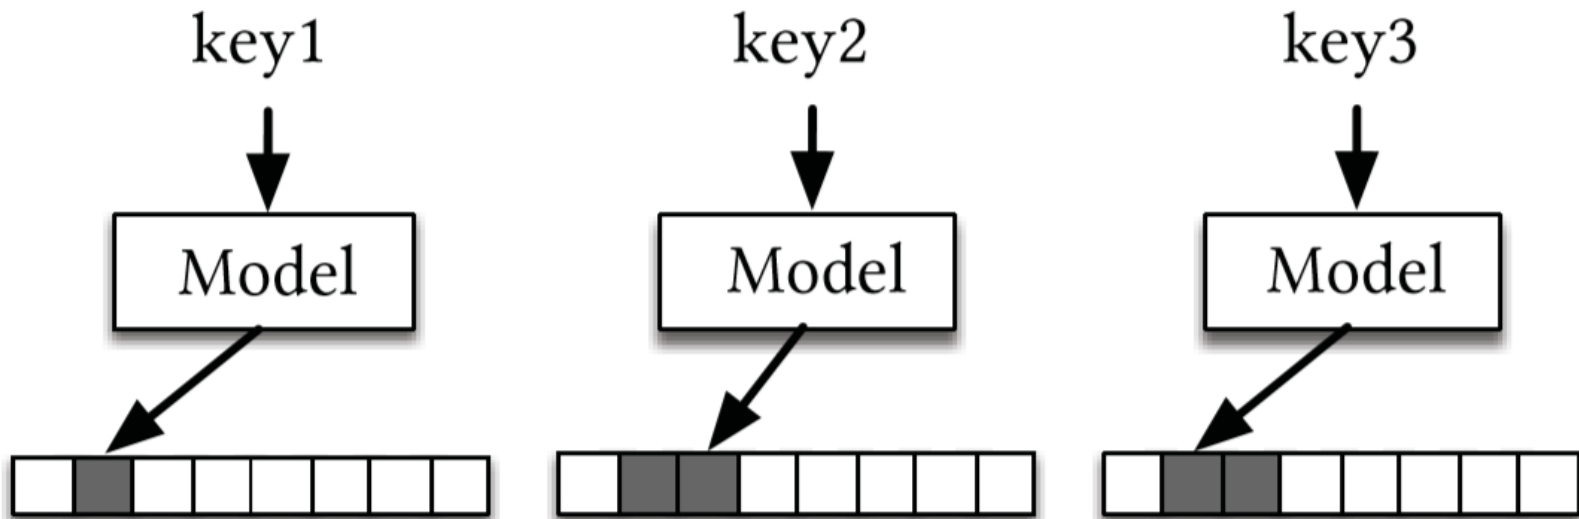
Learned Bloom filters as a classification problem



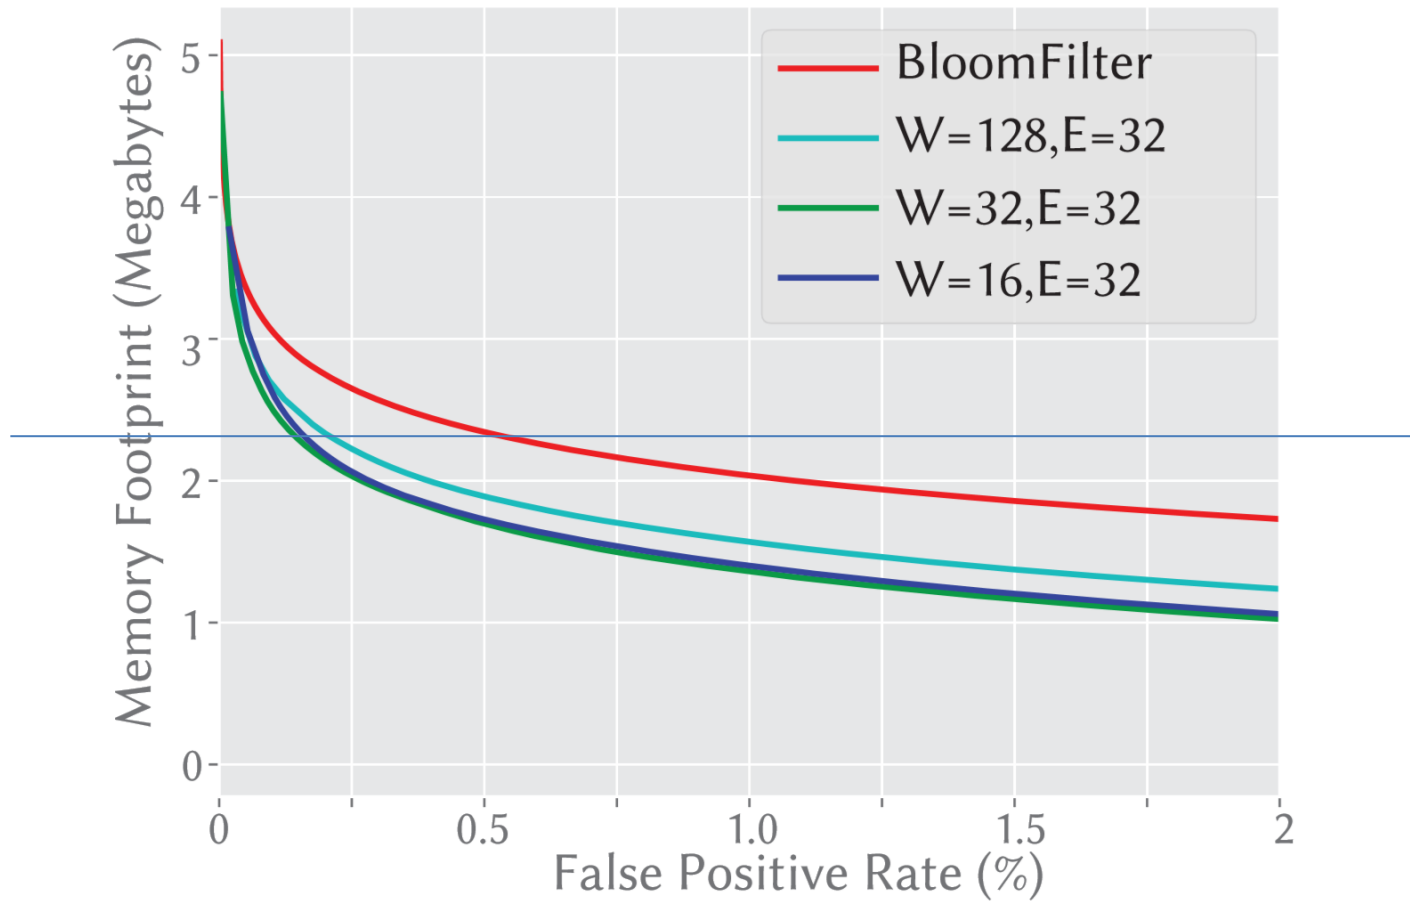
Goal: reduce False
Positive Rate

Learned Bloom filters with Model-Hashes

Map most keys to higher range of bit positions
and non-keys to the lower range.



Results



W: neuron number per layer of RNN
E: embedding size for each character

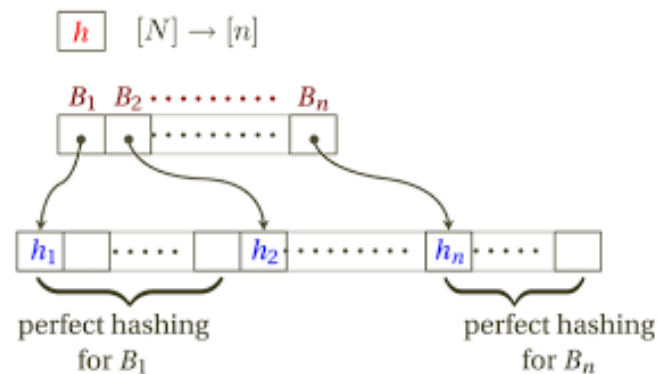
<https://huangliu0909.github.io>



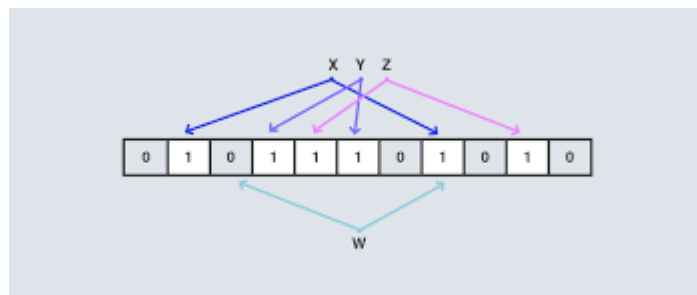
Related Work & Conclusion



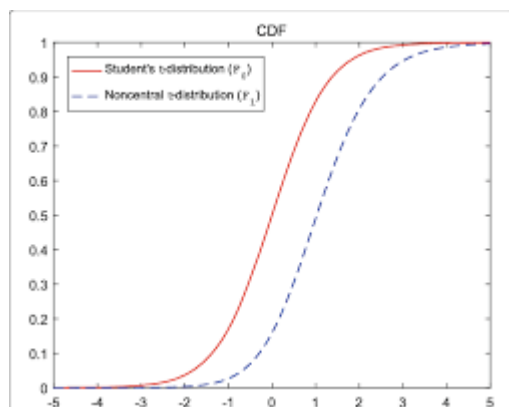
<https://images.app.goo.gl/5vU6ZDVxZYfLAChh9>



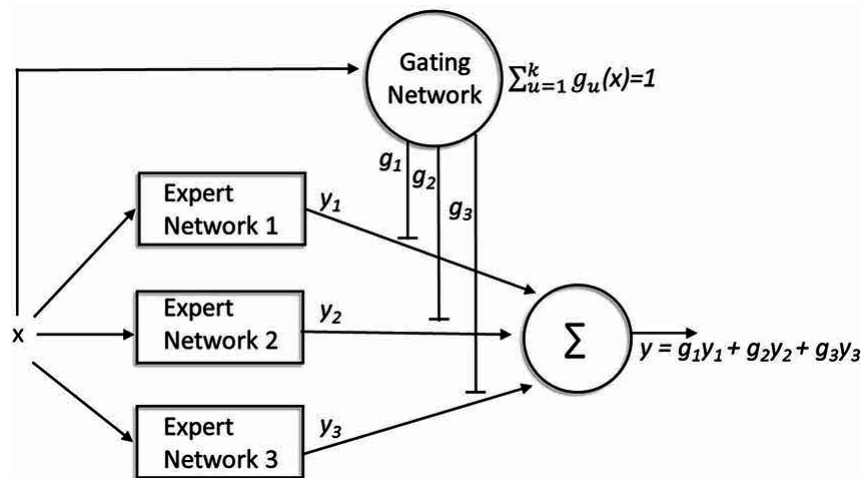
http://tcs.nju.edu.cn/wiki/index.php/%E9%9A%8F%E6%9C%BA%E7%AE%97%E6%B3%95_%28Fall_2011%29/Perfect_hashing



<https://images.app.goo.gl/3FFo7W8xXoD87XYdA>



<https://images.app.goo.gl/zZpd8BnccVutcMx66>



<https://images.app.goo.gl/m199gdAeoX94X7Xv8>

Future Work

- **Other ML Models**
- **Multi-Dimensional Indexes**
- **Learned Algorithms beyond indexing**
- **GPU/TPU**

“In summary, we have demonstrated that machine learned models have the potential to provide significant benefits over state-of-the-art indexes, and we believe this is a fruitful direction for future research.”

THANKS